



# Predicting New User Bookings

Anaëlia Ovalle, Michael Liston, Brent Rucker



# Table of contents

I. Introduction to the problem and our goal

II. Data exploration

III. Data pre-processing

IV. Approaches

# Introduction

New users on Airbnb can book a place to stay in 34,000+ cities across 190+ countries. By accurately predicting where a new user will book their first travel experience, Airbnb can share more personalized content with their community, decrease the average time to first booking, and better forecast demand.

The challenge is to predict in which country a new user will make his or her first booking.

Data available:

test\_users.csv and train\_users.csv : contains basic information and destination of past users.

Sessions.csv : contains additional web log informations generated when user use Airbnb website.

# Data exploration


Check for NULL values in features

```
[18] users_nan = (df.isnull().sum() / df.shape[0]) * 100
      users_nan[users_nan > 0].drop('country_destination')
```


```
↳ age          42.412365
   date_first_booking  67.733998
   first_affiliate_tracked  2.208335
   gender        46.990169
   language      0.000363
```

We have quite a lot of NaN in the age and gender which will yield in lesser performance of the classifiers we will build. The feature date\_first\_booking has a 67% of NaN values because this feature is not present at the test users, and therefore, we won't need it at the modeling part.

Distribution of data over destination country(target variable):



# Outliers in numerical variables:



Here is a boxplot of the age variable showing that there are some outliers far away from the logical age interval.

# Data pre-processing




- **Processing date columns.**

For date columns (`date_account_created`, `date_first_active`) we separate them to new columns that express meaningful time values like days of the week, month, and day of the month. we do that as these new columns should reflect specific habits or user behavior in specific times at the week or month.

- **Handling outliers.**

We remove age values that are greater than 100 or less than 16



Here is a boxplot showing age distribution after removing outliers




- **Missing values.**

For numerical variables, we fill any missing value with the mean of its variable. and for categorical variables, we fill it with the most frequent category.

another technique could be used with numerical columns, is to predict those missing values of the variable using other variables that are highly correlated with it.

So we draw a heatmap showing correlation of the variables with each other.

Pearson Correlation of Features



The heatmap showed that 'signup\_flow' is the most correlated variable to 'age'.

drawing the relation between 'signup\_flow' and 'age'.



It's showing a non-linear relation that can be captured by a decision tree or random forest regressor.

## Handling Imbalanced Data.

imbalanced data may lead to unstable training and low score when dealing with scoring metrics such as precision, recall or NDCG.

we can handle this issue with either under-sampling to the most represented classes or oversampling to the least ones.

- under-sampling could be done randomly by removing random samples, or using algorithms like Tomek links or Cluster Centroids.
- oversampling could be done by repeating minority class samples or using an algorithm like SMOTE

As we saw earlier that our data is imbalanced, So, I have used Tomeklinks to under-sample majority classes: 'NDF', 'US'. and SMOTE to oversample minority classes

# Convert categories to numbers

The next step is to convert all categories variables to numerical values so we can input to our models



first_affiliate_tracked	first_browser	first_device_type	gender	language	signup_app
untracked	Chrome	Mac Desktop	FEMALE	en	Web
untracked	Chrome	Mac Desktop	MALE	en	Web
untracked	IE	Windows Desktop	FEMALE	en	Web
untracked	Firefox	Mac Desktop	FEMALE	en	Web
untracked	Chrome	Mac Desktop	FEMALE	en	Web

first_affiliate_tracked	first_browser	first_device_type	gender	language	signup_app
6	8	3	0	5	2
6	8	3	1	5	2
6	22	6	0	5	2
6	17	3	0	5	2
6	8	3	0	5	2

# Approaches

I 've implemented three approaches to solve this problem

- XGBoost
- Random forest classifier
- Neural network

## **XGBoost:**

XGBoost is a decision-tree-based ensemble Machine Learning algorithm that uses a gradient boosting framework.

- I've implemented an XGBoost model with these hyperparameters.
- It got NDCG score = 0.837 on training data

```
from sklearn.metrics import ndcg_score
ndcg_score(true_relevance, y_pred)
```

0.8376089698333912

- And got NDCG = 0.852 on kaggle submission

[prediction.csv](#)

4 days ago by [mostafa shahin](#)

xgboost

```
params = {'max_depth': 10,
          'learning_rate': 1,
          'n_estimators': 5,
          'objective': 'multi:softprob',
          'num_class': 12,
          'gamma': 0,
          'min_child_weight': 1,
          'max_delta_step': 0,
          'subsample': 1,
          'colsample_bytree': 1,
          'colsample_bylevel': 1,
          'reg_alpha': 0,
          'reg_lambda': 1,
          'scale_pos_weight': 1,
          'base_score': 0.5,
          'missing': None,
          'nthread': 4,
          'seed': 42
        }
num_boost_round = 5
```

0.85262

0.84688



## Random forest classifier

I've implemented an rf classifier with 100 trees.

```
▶ print('training score : ' , scoring_fn(model,X_new_train,y_new_train))  
  print('validation score :', scoring_fn(model,X_valid,y_valid))
```

```
↳ training score : 0.9795342324160299  
  validation score : 0.8253954751378056
```

it did well on training-data, but it's not the case for validation-data.



then I set 'minimum samples per leaf' to 3

```
model = RandomForestClassifier(n_estimators=100, min_samples_leaf=3, n_jobs=-1)
```

validation score had slightly increased.

```
[ ] print('training score :', scoring_fn(model,X_new_train,y_new_train))  
    print('validation score :', scoring_fn(model,X_valid,y_valid))
```

```
↳ training score : 0.8957879921790012  
   validation score : 0.8362531356194036
```

Kaggle submission NDCG = 0.851

[prediction4.csv](#)

2 days ago by [mostafa shahin](#)

rf2

0.85169

0.84720

# Neural network

- usually, a neural network approach is mostly used for unstructured data problems but with the help of 'Entity Embedding' concept, we can use it for structured data. So, we can benefit from auto feature extraction that the neural network can introduce.
- I've implemented a neural network with a number of embedding layers(one for each categorical variable) prior to network hidden layers.

a simple network that consists of the embedding layers and three hidden layers got an NDCG score = 0.865 (higher than the previous two approaches)



```
neural_network3(  
  (embedd_layers): ModuleList(  
    (0): Embedding(8, 4)  
    (1): Embedding(18, 9)  
    (2): Embedding(7, 4)  
    (3): Embedding(55, 28)  
    (4): Embedding(9, 5)  
    (5): Embedding(3, 2)  
    (6): Embedding(25, 13)  
    (7): Embedding(4, 2)  
    (8): Embedding(18, 9)  
    (9): Embedding(4, 2)  
    (10): Embedding(7, 4)  
    (11): Embedding(31, 16)  
    (12): Embedding(12, 6)  
    (13): Embedding(7, 4)  
    (14): Embedding(31, 16)  
    (15): Embedding(12, 6)  
  )  
  (bn1): BatchNorm1d(138, eps=1e-05, momentum=0.1, affine=True, track_running_stats=True)  
  (fc1): Linear(in_features=138, out_features=64, bias=True)  
  (fc2): Linear(in_features=64, out_features=64, bias=True)  
  (fc3): Linear(in_features=64, out_features=12, bias=True)  
  (dropout): Dropout(p=0.25, inplace=False)  
)
```

prediction5.csv

30 minutes ago by mostafa shahin

neural network model

0.86535

0.86097